

INFO 251: Applied Machine Learning

Gradient Descent

A pen and paper might be useful for today's lecture (to sketch graphs)

Course Outline

- Causal Inference and Research Design
 - Experimental methods
 - Non-experiment methods
- **Machine Learning**
 - Design of Machine Learning Experiments
 - **Linear Models and Gradient Descent**
 - Non-linear models
 - Fairness and Bias in ML
 - Neural models
 - Deep Learning
 - Practicalities
 - Unsupervised Learning
- Special topics

Key Concepts (previous lecture)

- Decision boundaries
- Voronoi diagrams
- (K -)Nearest Neighbors
- Similarity and Distance metrics
- Normalization and Standardization
- Feature weighting

Key Concepts (today's lecture)


- Cost Functions
- Gradient Descent
- Local and global minima
- Convex functions
- Incremental vs. Batch GD
- Learning rates
- Feature scaling

Cost minimization

- In general:
 - We make a prediction of Y using some function $f(X)$
 - To choose the best model:
 - Define a loss function $J(Y, f(X))$
 - Minimize the expected loss of J
- With linear regression:
 - f is a linear function (e.g., $\alpha + \beta X$)
 - OLS regression minimizes squared-error loss $E(Y - f(X))^2$

Linear Regression

- OLS as Maximum Likelihood Estimation:
 - $Y_i = \alpha + \beta X_i + \epsilon_i$
 - Idea: Choose α and β so that $\alpha + \beta X_i$ is “as close as possible” to Y_i for training data
- In other words
 - $\min_{\alpha, \beta} \sum_{i=1}^N (\alpha + \beta X_i - Y_i)^2$
- In general, we are minimizing a **Cost Function J**
 - $\min_{\alpha, \beta} J(\alpha, \beta)$
 - In the case of OLS, we use a “squared error” cost function
 - $J(\alpha, \beta) = \frac{1}{2N} \sum_{i=1}^N (\alpha + \beta X_i - Y_i)^2$



General formulation (OLS)

- Model ("hypothesis")

- $Y_i = \alpha + \beta X_i$

- Parameters


- α, β

- Cost Function

- $J(\alpha, \beta) = \frac{1}{2N} \sum_{i=1}^N (\alpha + \beta X_i - Y_i)^2$

- Objective

- $\min_{\alpha, \beta} J(\alpha, \beta)$



OLS with no intercept

- Model ("hypothesis")

- $Y_i = \beta X_i$

- Parameters


- β

- Cost Function

- $J(\beta) = \frac{1}{2N} \sum_{i=1}^N (\beta X_i - Y_i)^2$


- Objective

- $\min_{\beta} J(\beta)$




Fill in the blanks

- Model (“hypothesis”)
 - Income is a linear function of education and also education², i.e. nonlinearities exist
 - $Y_i = \alpha + \beta X_i + \gamma X_i^2$
- Parameters
 - α, β, γ
- Cost Function
 - Use “absolute error” cost function
 - $J(\alpha, \beta, \gamma) = \frac{1}{N} \sum_{i=1}^N |\alpha + \beta X_i + \gamma X_i^2 - Y_i|$
- Objective
 - $\min_{\alpha, \beta, \gamma} J(\alpha, \beta, \gamma)$





Exercise: Computing Cost

- Assume our data look like this ($N = 3$)
 - $X_1 = 1, Y_1 = 1$
 - $X_2 = 2, Y_2 = 2$
 - $X_3 = 3, Y_3 = 3$
- Our model is $Y_i = \beta X_i$
 - Our cost function is squared error: $J(\beta) = \sum_{i=1}^N (\beta X_i - Y_i)^2$
- Your task is to compute $J(\beta)$, given these 3 points, for:
 - $\beta = 1$
 - $\beta = 0$
 - $\beta = 2$
 - $\beta = 0.5$ (you might need a calculator)
- Draw a plot of $J(\beta)$ as a function of β




Visualizing Cost (1 parameter)

- Where is $J(\beta)$ minimized?



Visualizing Cost (2 parameters)

- Generalizing to a multi-dimensional loss surface
 - Model ("hypothesis"): $Y_i = \theta_1 + \theta_2 X_i$
 - Objective: $\min_{\theta_1, \theta_2} J(\theta_1, \theta_2)$




Outline

- Cost functions
- **Gradient descent**
- Feature scaling


Gradient Descent

- Gradient descent provides a principled method/algorithm to minimize the cost function J
- Idea: to solve $\min_{\alpha, \beta} J(\alpha, \beta)$
 - Initialize α, β
 - Change α, β in some way that reduces $J(\alpha, \beta)$
 - Eventually we will end up at a minimum
- What about analytic solution: $(X'X)^{-1}X'Y$
 - Sometimes not practically feasible (too much data, multicollinearity)

Gradient Descent: Visualization



Local Minima



Gradient Descent Algorithm (incremental)

- In pseudo-code:

Choose an initial vector of parameters α, β

Choose learning rate R

Repeat until convergence (i.e., until an approximate minimum is obtained):

For each example i in training set:


$$\left. \begin{aligned} \alpha &\leftarrow \alpha - R \frac{\partial}{\partial \alpha} J(\alpha, \beta) \\ \beta &\leftarrow \beta - R \frac{\partial}{\partial \beta} J(\alpha, \beta) \end{aligned} \right\} \text{Simultaneous update}$$

- With multiple predictors/regressors...

- $Y_i = \alpha + \beta_1 X_{i1} + \beta_2 X_{i2} + \dots + \beta_k X_{ik}$


Gradient Descent: Derivative

- In 1-Dimension
- Update Rule:
 - $\beta \leftarrow \beta - \mathbb{R} \frac{\partial}{\partial \beta} J(\alpha, \beta)$
- Initialize β at 1.9
 - What's the derivative $\frac{\partial}{\partial \beta} J(\alpha, \beta)$?
 - How does β update?
- Initialize β at 0.2
 - What's the derivative $\frac{\partial}{\partial \beta} J(\alpha, \beta)$?
 - How does β update?



Gradient Descent: Learning Rate

- $\beta \leftarrow \beta - R \frac{\partial}{\partial \beta} J(\alpha, \beta)$
- What does R do?
- Small R:
 - Gradient descent can be slow
- Large R:
 - Can overshoot the minimum
 - May fail to converge
 - May diverge!



Gradient Descent: Convergence

- Do we need to change the learning rate?

Choose an initial vector of parameters α, β

Choose learning rate R

Repeat until convergence:

For each example i :

$$\alpha \leftarrow \alpha - R \frac{\partial}{\partial \alpha} J(\alpha, \beta)$$

$$\beta \leftarrow \beta - R \frac{\partial}{\partial \beta} J(\alpha, \beta)$$

- Not typically. Gradient descent can converge to a local minimum, even with the learning rate fixed
 - As we approach a local minimum, gradient descent takes smaller steps
 - But adaptive learning rates can help speed up convergence, prevent overshooting

Gradient Descent: Regression

- Gradient Descent

Repeat until convergence:

$$\alpha \leftarrow \alpha - \text{R} \frac{\partial}{\partial \alpha} J(\alpha, \beta)$$

$$\beta \leftarrow \beta - \text{R} \frac{\partial}{\partial \beta} J(\alpha, \beta)$$

- Regression cost function

- $J(\alpha, \beta) = \frac{1}{2N} \sum_{i=1}^N (\alpha + \beta X_i - Y_i)^2$

- The missing pieces: $\frac{\partial}{\partial \alpha} J(\alpha, \beta)$ and $\frac{\partial}{\partial \beta} J(\alpha, \beta)$

- $\frac{\partial}{\partial \alpha} J(\alpha, \beta) = \frac{\partial}{\partial \alpha} \frac{1}{2N} \sum_{i=1}^N (\alpha + \beta X_i - Y_i)^2$

- $\frac{\partial}{\partial \beta} J(\alpha, \beta) = \frac{\partial}{\partial \beta} \frac{1}{2N} \sum_{i=1}^N (\alpha + \beta X_i - Y_i)^2$

Gradient Descent: Regression

- Partial derivatives:

$$\begin{aligned}
 \frac{\partial}{\partial \alpha} J(\alpha, \beta) &= \frac{\partial}{\partial \alpha} \frac{1}{2N} \sum_{i=1}^N (\hat{Y}_i - Y_i)^2 \\
 &= \frac{\partial}{\partial \alpha} \frac{1}{2N} \sum_{i=1}^N (\alpha + \beta X_i - Y_i)^2 \\
 &= \frac{1}{2N} \sum_{i=1}^N \frac{\partial}{\partial \alpha} (\alpha + \beta X_i - Y_i)^2 \\
 &= \frac{1}{2N} \sum_{i=1}^N 2(\alpha + \beta X_i - Y_i) \frac{\partial}{\partial \alpha} (\alpha + \beta X_i - Y_i) \\
 &= \frac{1}{N} \sum_{i=1}^N (\alpha + \beta X_i - Y_i) \\
 &= \frac{1}{N} \sum_{i=1}^N (\hat{Y}_i - Y_i)
 \end{aligned}$$

$$\begin{aligned}
 \frac{\partial}{\partial \beta} J(\alpha, \beta) &= \frac{\partial}{\partial \beta} \frac{1}{2N} \sum_{i=1}^N (\hat{Y}_i - Y_i)^2 \\
 &= \frac{1}{2N} \sum_{i=1}^N \frac{\partial}{\partial \beta} (\alpha + \beta X_i - Y_i)^2 \\
 &= \frac{1}{2N} \sum_{i=1}^N 2(\alpha + \beta X_i - Y_i) \frac{\partial}{\partial \beta} (\alpha + \beta X_i - Y_i) \\
 &= \frac{1}{N} \sum_{i=1}^N (\alpha + \beta X_i - Y_i) (X_i) \\
 &= \frac{1}{N} \sum_{i=1}^N (\hat{Y}_i - Y_i) X_i
 \end{aligned}$$

Gradient Descent: Regression


- Gradient Descent Algorithm (linear regression)

Repeat until convergence:


$$\alpha \leftarrow \alpha - \mathbb{R} \frac{1}{N} \sum_{i=1}^N (\alpha + \beta X_i - Y_i)$$

$$\beta \leftarrow \beta - \mathbb{R} \frac{1}{N} \sum_{i=1}^N (\alpha + \beta X_i - Y_i) X_i$$


Gradient Descent: In Action




Gradient Descent: In Action




Gradient Descent: In Action




Gradient Descent: In Action




Gradient Descent: In Action




Gradient Descent: In Action




Gradient Descent: In Action



Gradient Descent: In Action



Gradient Descent: In Action

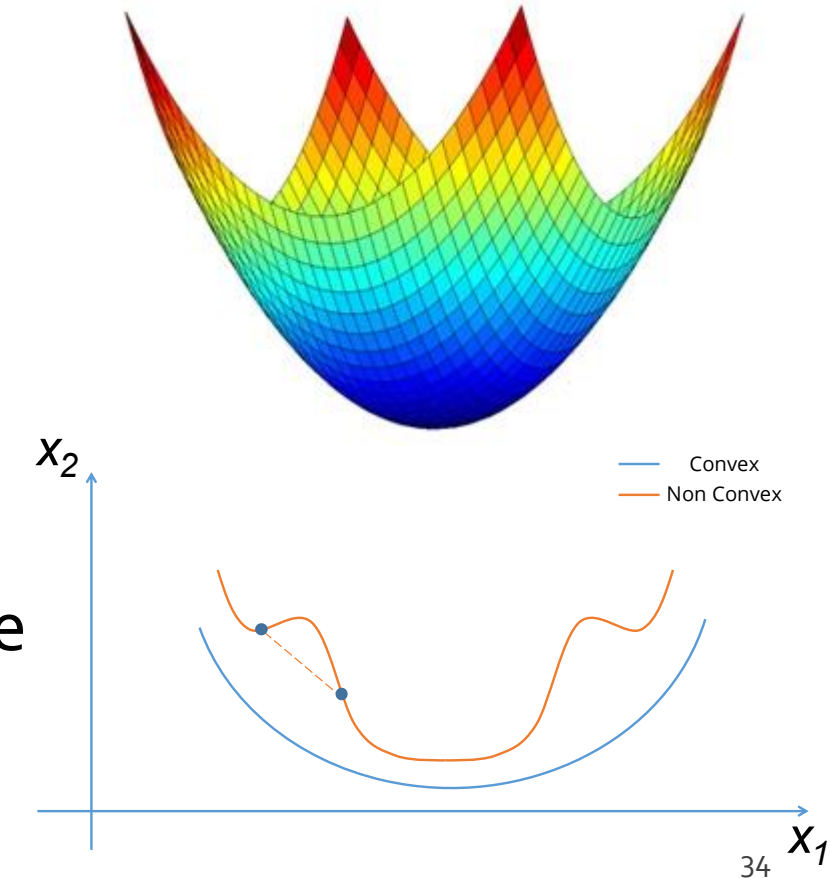



Outline

- Cost functions
- Gradient descent
 - **Local Minima**
 - Batch and Incremental versions
- Feature scaling


Local Minima?

- What about local minima in gradient descent for regression?
- No problem!
- Cost function in regression is **convex**
 - Convex: second derivative is non-negative
 - More intuitively: a continuous function where the midpoint of any interval doesn't exceed the mean of the endpoints



Local Minima?

- What if cost function is not convex?
- Several options
 - Use multiple initialization points
 - Or use “smart” starting points (e.g., Xavier, He initialization)
 - Momentum can help
 - E.g., Nesterov Accelerated Gradient (NAG), RMSprop, Adam
 - Adaptive learning rates can help



Stopping conditions

Choose an initial vector of parameters α, β

Choose learning rate R

Repeat until convergence (i.e., until an approximate minimum is obtained):

$$\alpha \leftarrow \alpha - R \frac{\partial}{\partial \alpha} J(\alpha, \beta)$$

$$\beta \leftarrow \beta - R \frac{\partial}{\partial \beta} J(\alpha, \beta)$$

- How to know a minimum has been obtained?
 - Look for small changes in the gradient
 - Look for small improvements in cost
 - Look for no changes in parameters
 - Set a stopping condition!

Outline

- Cost functions
- Gradient descent
 - Local Minima
 - **Batch and Incremental versions**
 - Other issues
- Feature scaling

Incremental vs. Batch Gradient Descent

- In “Batch” gradient descent


Repeat until convergence:

$$\begin{aligned} \text{Compute } \nabla \alpha &= \frac{\partial}{\partial \alpha} J(\alpha, \beta) = \frac{1}{N} \sum_{i=1}^N (\alpha + \beta X_i - Y_i) \\ \text{Compute } \nabla \beta &= \frac{\partial}{\partial \beta} J(\alpha, \beta) = \frac{1}{N} \sum_{i=1}^N (\alpha + \beta X_i - Y_i) X_i \end{aligned} \quad \left. \vphantom{\begin{aligned} \text{Compute } \nabla \alpha \\ \text{Compute } \nabla \beta \end{aligned}} \right\} \text{Global gradient, computed over all training data}$$

$$\begin{aligned} \alpha &\leftarrow \alpha - R \nabla \alpha \\ \beta &\leftarrow \beta - R \nabla \beta \end{aligned} \quad \left. \vphantom{\begin{aligned} \alpha &\leftarrow \alpha - R \nabla \alpha \\ \beta &\leftarrow \beta - R \nabla \beta \end{aligned}} \right\} \text{Single, simultaneous update}$$

- Note: each step uses all training examples!

Batch Gradient Descent



Incremental vs. Batch Gradient Descent

- “Iterative” (stochastic) version of gradient descent:

Repeat until an approximate minimum is obtained:

Randomly shuffle examples in the training set


For each example i :

$$\alpha \leftarrow \alpha - \mathbb{E} \frac{\partial}{\partial \alpha} J(\alpha, \beta) \quad // \text{ evaluate } \frac{\partial}{\partial \alpha} J(\alpha, \beta) \text{ at } x_i \text{ and update } \alpha$$

$$\beta \leftarrow \beta - \mathbb{E} \frac{\partial}{\partial \beta} J(\alpha, \beta) \quad // \text{ evaluate } \frac{\partial}{\partial \beta} J(\alpha, \beta) \text{ at } x_i \text{ and update } \beta$$

- The parameters are adjusted with each training instance, iteratively
- Also: Mini-batch

Stochastic Gradient Descent





Outline

- Cost functions
- Gradient descent
- **Feature scaling**

Feature scaling

- In gradient descent, we “take a step” in the direction where the decrease in cost is greatest
- When some features (axes) are on different scales, gradient descent can be inefficient
 - Putting different features on same scale can make gradient descent much faster



Feature scaling

- Feature scaling is an important pre-processing step for many common machine learning algorithms

- Standardization: (mean 0 standard dev. 1)

$$x'_{ik} = \frac{x_{ik} - \bar{x}_k}{s_k}$$

- s_k is typically standard deviation of x_k or range (max-min) of x_k

- Also common: force feature to be roughly between -1 and 1:

$$x'_{ik} = \frac{x_{ik}}{\max(|x_k|)}$$

- Note: normalization parameters should be learned **on training data**

Key Concepts (today's lecture)

- Cost Functions
- Gradient Descent
- Local and global minima
- Convex functions
- Incremental vs. Batch GD
- Learning rates
- Feature scaling