



Thanks to machine-learning algorithms,  
the robot apocalypse was short-lived.

INFO 251: Applied Machine Learning

# Random Forests

# Key Concepts (Decision Trees)

- Churn prediction
- Decision tree representation
- Hyper-rectangles and decision boundary
- Recursive tree building algorithm
- Splitting
- Information gain

# Intuition check

- True or false:
  - The recursive decision tree algorithm is deterministic (assuming we are not performing cross-validation and have fixed the hyperparameters in advance)

# Course Outline

- Causal Inference and Research Design
  - Experimental methods
  - Non-experiment methods
- **Machine Learning**
  - Design of Machine Learning Experiments
  - Linear Models and Gradient Descent
  - Fairness and Bias in ML
  - **Non-linear models**
  - Neural models
  - Deep Learning
  - Practicalities
  - Unsupervised Learning
- Special topics

# Outline

- Regression Trees
- Random Forests
- Ensembles
- Boosting
- Feature Importance

# Key Concepts (Random Forests)

- Regression vs. Decision trees
- Recursive regression trees algorithm
- Random forests
- Bagging
- Stacking
- Adaboost
- Gradient boosting
- Feature importance

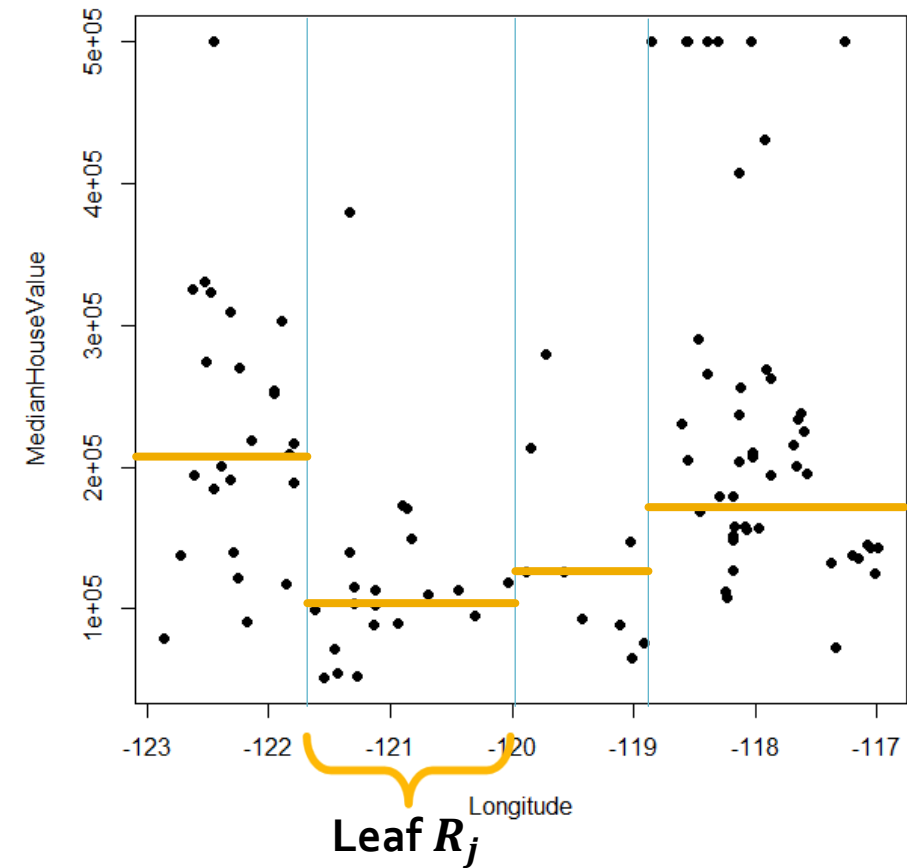
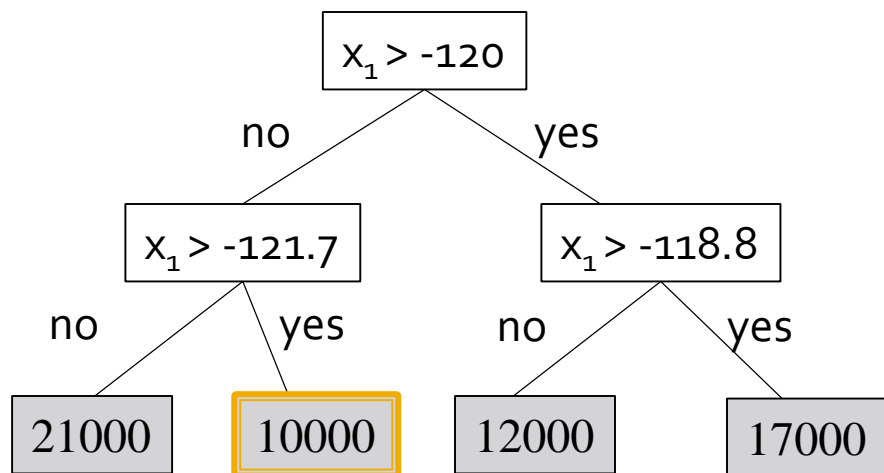
# Regression Trees

- What if output values are continuous or real-valued (i.e., not discrete)?
- Regression trees
  - Construct binary tree, minimize error in each leaf
  - Before, we counted # elements of each type in leaf
  - Now we choose predicted value that minimizes error
- Example: Predict median housing value based on a house's location (latitude, longitude)

# Regression Trees

```
> head(calif[,c(1,8,9)])
```


	MedianHouseValue	Latitude	Longitude
1	452600	37.88	-122.23
2	358500	37.86	-122.22
3	352100	37.85	-122.24
4	341300	37.85	-122.25
5	342200	37.85	-122.25
6	269700	37.85	-122.25
	...		





# Regression Trees


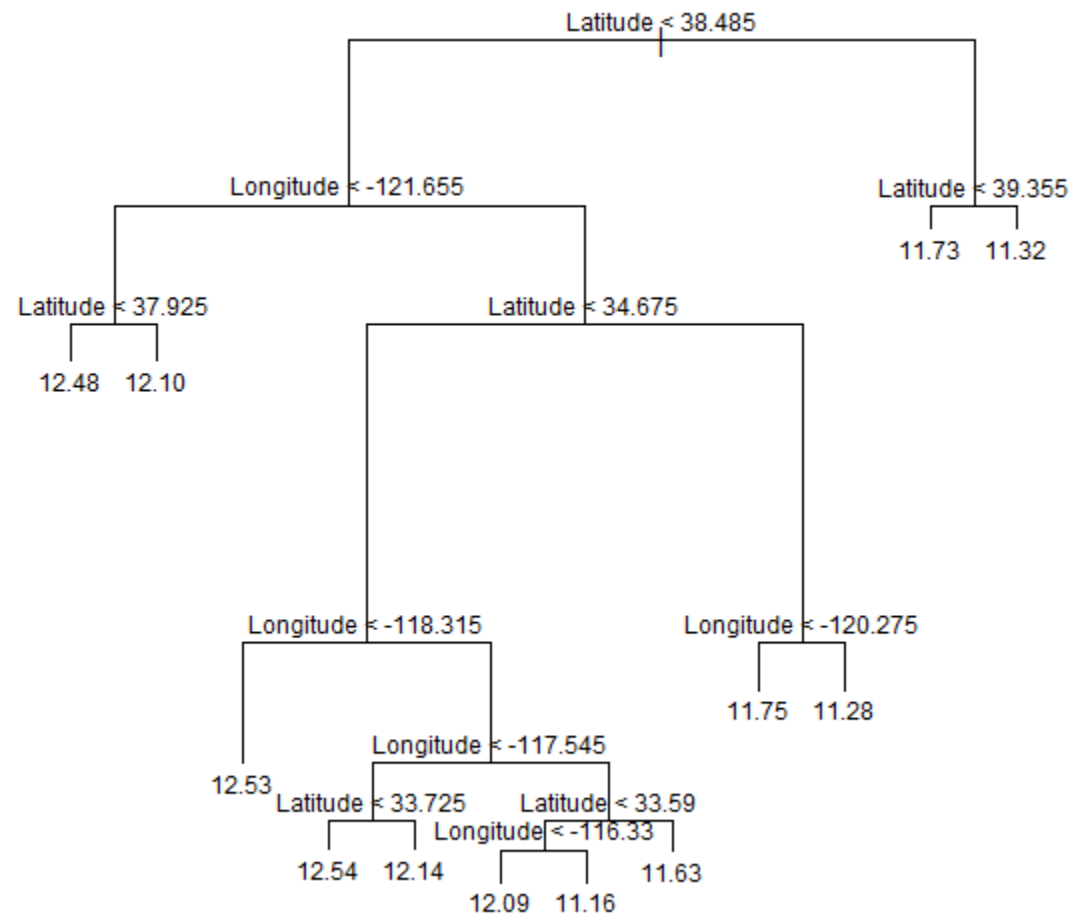
- How to choose split point?
  - Idea: Minimize prediction error
- In 1-dimension: choose  $s$  to minimize
  - $\min_{\bar{y}} \sum_{i: x_i > s} (\bar{y} - y_i)^2 + \min_{\bar{y}'} \sum_{i: x_i \leq s} (\bar{y}' - y_i)^2$
  - Consider finite splits (e.g.  $s$  between data)
- This intuition generalizes to  $D$  dimensions



# Regression Trees: Recursive Algorithm

1. Start with a single node ( $c_j$ ) containing all points.
  1. Calculate predicted value  $\bar{y}_{c_j} = \frac{1}{n} \sum_{i \in c_j} y_i$
  2. Calculate total error:  $J = \sum_{c_j} \sum_{i \in c_j} (\bar{y}_{c_j} - y_i)^2$
2. If all points in the node have identical features (predictors), stop.
  - Otherwise, search all binary splits of all variables for split that most reduces  $J$
  - Stop if  $J$  decreases less than  $\delta$  or if nodes are close to empty
  - Otherwise, make that split, creating two new nodes
3. Recurse on each new node

# Regression Trees: Example




# Outline

- Regression Trees
- **Random Forests**
- Ensembles
- Boosting
- Feature Importance

# Trees to forests

- Which classifier works best?
  - “Random forests” combine outputs of multiple classifiers



# Building a forest

- Bootstrap sample a new training set
  - with replacement
- Build a decision tree
  - The randomization (of the sample and/or features) forces differentiated trees
  - Optional: randomly select a subset of features
  - Pruning not required! “Regularization” occurs through forest
- Repeat until you have lots of trees
- Predict by taking a vote among the trees

# Example: The “CART” forest


- Formally:

$$\hat{y}_i = \sum_{f_k \in \mathcal{F}} f_k(x_i)$$

- $\mathcal{F}$  is the space of regression trees
  - Each  $f_k$  maps data examples  $x_i$  to tree leaves
  - Scores are averaged (or summed, depending on implementation) across trees

# Other ensemble methods

- Bagging = **bootstrap aggregating**
  - Create artificial versions of data via bootstrap
  - 1 sample = bootstrap
  - $M$  samples = bagging
- Stacking: train model (e.g. another tree, a logistic regression) on output of other models
- Boosting (Kearns, 1988)
  - Train a sequence of models, each emphasizes the examples misclassified by the previous model



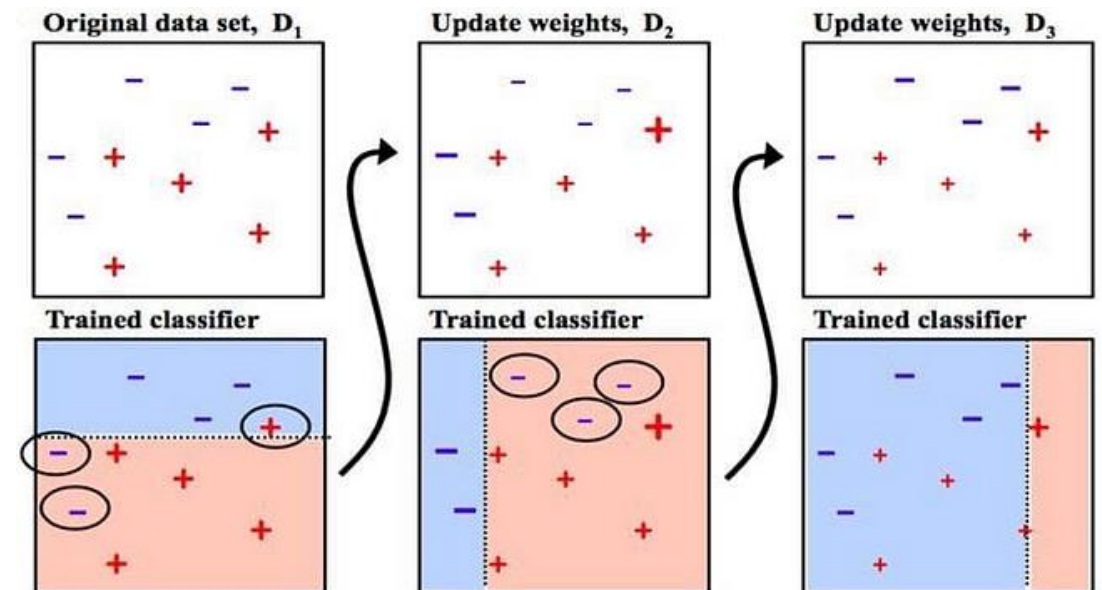


# Outline

- Regression Trees
- Random Forests
- Ensembles
- **Boosting**
- Feature Importance

# Adaptive Boosting (AdaBoost)

- Adaboost:
  1. Initially, set a uniform weight for each training example =  $1/n$
  2. Train a classifier where the objective respects the weights
    - Could be any classifier, but original Adaboost used single-feature decision stump
  3. Increase the weights for misclassified examples
  4. Return to 2



# Gradient Boosting



Linear Regression



Gradient Boosting

# (Extreme) Gradient Boosting

- Start with regression tree-based model:

$$\hat{y}_i = \sum_{f_k \in \mathcal{F}} f_k(x_i)$$

- Gradient boosting loss function “fits on residuals”:

$$\mathcal{L}^{(t)} = \sum_{i=1}^n l(y_i, \underbrace{\hat{y}_i^{(t-1)} + f_t(x_i)}_{f_t \text{ fits on residual of } t-1}) + \underbrace{\gamma T + \frac{1}{2} \lambda \|w\|^2}_{\text{Regularization penalties}}$$

- $T$  is the number of leaves
  - $t$  indexes training iterations
  - $w$  is vector of scores on each leaf (i.e., the leaf weights)
- Optimization is similar to gradient descent
  - Relies on being able to measure how good each tree is
  - Next tree solves for the loss of prior tree

**Details** Friedman, J.H., 2001. Greedy function approximation: a gradient boosting machine. Annals of statistics 1189–1232.

**See also:** Chapter 10.9 and 10.10 of Elements of Statistical Learning

# Outline

- Regression Trees
- Random Forests
- Ensembles
- Boosting
- **Feature Importance**


# Feature Importance

- The primary focus of random forests is prediction, rather than inference
  - A single tree is fairly interpretable, but it's hard to interpret a forest
  - This is generally true for complex, non-parametric, and non-linear models
- Nonetheless, people frequently still want to do some ex-post interpretation
  - "What features are important to the classifier?"

# Feature Importance


- Intuitively, the features that “matter”:
  - Occur high in the tree (high information gain for that tree)
  - Occur frequently in the tree (if feature is non-binary)
  - Occur in many trees (if it's a forest)

# Feature Importance: Example





# Feature Importance: Example



# Feature Importance

- Formally, two common approaches:
  1. **Mean decrease impurity (aka Gini importance):** average (across trees) decrease in weighted impurity caused by that feature

```

1  from sklearn.datasets import load_boston
2  from sklearn.ensemble import RandomForestRegressor
3  import numpy as np
4  #Load boston housing dataset as an example
5  boston = load_boston()
6  X = boston["data"]
7  Y = boston["target"]
8  names = boston["feature_names"]
9  rf = RandomForestRegressor()
10 rf.fit(X, Y)
11 print "Features sorted by their score:"
12 print sorted(zip(map(lambda x: round(x, 4), rf.feature_importances_), names),
13               reverse=True)

```

Features sorted by their score:

```

[(0.5298, 'LSTAT'), (0.4116, 'RM'), (0.0252, 'DIS'), (0.0172, 'CRIM'), (0.0065, 'NOX'),
(0.0035, 'PTRATIO'), (0.0021, 'TAX'), (0.0017, 'AGE'), (0.0012, 'B'), (0.0008, 'INDUS'),
(0.0004, 'RAD'), (0.0001, 'CHAS'), (0.0, 'ZN')]

```

# Feature Importance

- Issues with impurity:
  - Biased towards features with multiple values
  - What happens when two features are closely correlated?

```
1 size = 10000
2 np.random.seed(seed=10)
3 X_seed = np.random.normal(0, 1, size)
4 X0 = X_seed + np.random.normal(0, .1, size)
5 X1 = X_seed + np.random.normal(0, .1, size)
6 X2 = X_seed + np.random.normal(0, .1, size)
7 X = np.array([X0, X1, X2]).T
8 Y = X0 + X1 + X2
9
10 rf = RandomForestRegressor(n_estimators=20, max_features=2)
11 rf.fit(X, Y);
12 print "Scores for X0, X1, X2:", map(lambda x:round (x,3),
13                                     rf.feature_importances_)
```

Scores for X0, X1, X2: [0.278, 0.66, 0.062]

# Feature Importance

- Formally, two common approaches:
  1. **Mean decrease impurity**: average (across trees) decrease in weighted impurity caused by that feature
  2. **Mean decrease accuracy (“Permutation Importance”)**: average (across trees) decrease in performance when a given feature is randomized
    - Not implemented in sklearn (but very easy to do so by hand - see ESLII reading)

# Feature Importance

- Issues
  - Interpret feature importances at your own risk!
  - They are informative, but rather atheoretical

# Recap

- Regression
  - Parametric, fast training, linear
- Nearest Neighbors
  - Non-parametric, no training, complex decisions
- Naïve Bayes
  - Parametric, very fast training
- Decision Trees
  - Non-linear decisions, intuitive model

# Key Concepts (this lecture)

- Regression vs. Decision trees
- Recursive regression trees algorithm
- Random forests
- Bagging
- Stacking
- Adaboost
- Gradient boosting
- Feature importance

# For Next Class:

- Read:
  - Daume, chapters 4 and 10
- Keep working on Problem Set 4!

