

INFO 251: Applied Machine Learning

Foundations of Neural Networks

OH, HEY, YOU ORGANIZED
OUR PHOTO ARCHIVE!

YEAH, I TRAINED A NEURAL
NET TO SORT THE UNLABELED
PHOTOS INTO CATEGORIES.

WHOA! NICE WORK!



ENGINEERING TIP:

WHEN YOU DO A TASK BY HAND,
YOU CAN TECHNICALLY SAY YOU
TRAINED A NEURAL NET TO DO IT.

Course Outline

- Causal Inference and Research Design
 - Experimental methods
 - Non-experiment methods
- **Machine Learning**
 - Design of Machine Learning Experiments
 - Linear Models and Gradient Descent
 - Non-linear models
 - Fairness and Bias in ML
 - **Neural models**
 - Deep Learning
 - Practicalities
 - Unsupervised Learning
- Special topics

Key Concepts (Trees and Forests)

- Decision trees and regression trees
- Recursive tree algorithm
- Choosing splits
- Information gain
- Overfitting and pruning
- Regression trees
- Random forests
- AdaBoost
- Gradient boosting
- Feature Importance

Outline

- **Neural Networks: Motivation and Biology**
- The Perceptron
- Learning perceptron weights
- Multilayer networks
- Learning multilayer weights

Neural Networks: Common Applications

- Basic building block of many modern breakthroughs in AI
 - Computer vision: facial recognition, object detection
 - Natural language: chatbots, GPT, Siri, Google Assistant
 - Autonomous systems: self-driving cars, real-time decisions
 - Healthcare: Medical imaging analysis, drug discovery
 - Finance: Fraud detection, market forecasting
 - Etc., etc.

Neural Networks

- Computational models inspired by the brain
 - Early work dates back to 1940's (McCullough & Pitts, 1943)
 - Idea: mimic how the brain processes information
 - In the hopes that computers can reason as well as human brains
 - ...and perhaps even better!
- So, how do real neurons work? (cue expert)

What's a Neuron?



What's a Neuron?





Outline

- Neural Networks: Motivation and Biology
- **The Perceptron**
- Learning perceptron weights
- Multilayer networks
- Learning multilayer weights: Intuition
- Generalizing logistic regression



Creating Artificial Neurons

- How to model a neuron?
 - A real neuron fires when membrane potential exceeds a threshold
- The “Perceptron” (Rosenblatt, 1958)
 - Simple binary threshold function
 - Perceptron “fires” if weighted sum of inputs exceeds threshold
 - $h(x) = \text{Sign}(b + \sum_{d=1}^k w_d x_d)$
 - k weights indexed by w_d
 - Bias term b (or w_0) allows for non-zero threshold



Linearly separable data


- Simple model, but revolutionary at the time!
- Perceptron works with **linearly separable** data
 - i.e., boundary can be specified by hyperplane
 - E.g., $w_0 + w_1x_1 + \dots + w_kx_k = 0$
 - Example: what formula defines the separating hyperplane for these data?



Perceptron: Examples

- A perceptron for AND:
 - Two weights and intercept:
 - $h(x_i) = w_0 + w_1x_{i1} + w_2x_{i2}$
 - One solution:
 - $w_1=0.5, w_2=0.5, w_0=-0.8$
 - *Can you find another one?*

x_1	x_2	y
1	1	T
1	0	F
0	1	F
0	0	F




Note: You'll sometimes see a threshold T used instead of the bias w_0 , such that $T = -w_0$ (in this example, $T=0.8$)

Perceptron: Your turn

- A perceptron for OR:
 - Two weights and intercept:
 - $h(x_i) = w_0 + w_1x_{i1} + w_2x_{i2}$
 - Find possible weights w_0, w_1, w_2




x_1	x_2	y
1	1	T
1	0	T
0	1	T
0	0	F



Perceptron: Examples

- You've seen AND and OR
- A perceptron for XOR?
 - Impossible! (Minsky & Papert 1969) → Why?
 - XOR is not **linearly separable**

x_1	x_2	y
1	1	F
1	0	T
0	1	T
0	0	F




Outline

- Neural Networks: Motivation and Biology
- The Perceptron
- **Learning perceptron weights**
- Multilayer networks
- Learning multilayer weights: Intuition
- Generalizing logistic regression

Learning weights

- Given we have input and output (for instance, a truth table), how do we learn the weights?
- Early insight: Hebbian Learning and Synaptic Plasticity (Hebb, 1949)
 - “Neurons that fire together, wire together”
- Modern networks learn using a variety of ways
 - We’ll start with Rosenblatt’s algorithm (1958)



Learning weights (Rosenblatt)

- Rosenblatt's Algorithm (perceptron):

initialize weights randomly

while termination condition is not met:

 initialize $\Delta w_j = 0$

 for each training example (X_i, Y_i) :

 compute predicted output \hat{Y}_i


 foreach weight w_j :

$\Delta w_j = \Delta w_j + \eta (Y_i - \hat{Y}_i) X_i$  "error-driven" learning

 for each weight w_j :

$w_j = w_j + \Delta w_j$  Learning rate

Perceptron: In action



Who cares?

- Rosenblatt proved the algorithm is guaranteed to converge as long as:
 - Training data are linearly separable
 - Learning rate is sufficiently small
 - (In the proof, it has to be infinitesimally small)

Training Rule vs. Gradient Descent

- This looks a lot like gradient descent
- Are these approaches different?
 - Training Rule (Rosenblatt)
 - $\Delta w_j = \Delta w_j + \eta (Y_i - \hat{Y}_i) X_i$
 - Gradient Descent w/ Logistic Regression
 - $\beta \leftarrow \beta + R(Y_i - \hat{Y}_i) X_i$
- The key is the \hat{Y}_i
 - Perceptron: \hat{Y}_i is a step function, either 0 or 1
 - G.D. requires convex surface, not a step function
 - Logit: \hat{Y}_i is a smooth, continuous function

Training Rule vs. Gradient Descent

- Perceptron Training Rule
 - Guaranteed to work if data are linearly separable
 - Requires sufficiently small learning rate η
- Training with Gradient Descent
 - With convex loss...
 - Guaranteed to converge to minimum error
 - Works when data contains noise
 - Works when data are not linearly separable

Perceptron: Summary


- **Online** algorithm: only considers one instance at a time
- **Error-driven**: Only updates on failure
- Guaranteed to converge if solution exists
- But boundary is **linear**

Outline

- Neural Networks: Motivation and Biology
- The Perceptron
- Learning perceptron weights
- **Multilayer networks**
- Learning multilayer weights: Intuition
- Generalizing logistic regression

Limitations of the Perceptron



- Only works with linearly separable data
- Only works if learning rate is small enough (Rosenblatt's proof)
- These sort of problems led to “long winter” (1980's)



x_1	x_2	y
1	1	-1
1	0	1
0	1	1
0	0	-1

Multilayer Networks


- Single-layer networks are limited – they can only learn hyperplanes
 - Most real-world problems are more complex
- What if we layer neurons?
 - **Multi-Layer Perceptron (MLP)**: an input layer, one or more hidden layers, and an output layer
 - E.g., two-layer network (two layers of weights)
 - This allows for *very* powerful and complex computation!



Nonlinearity

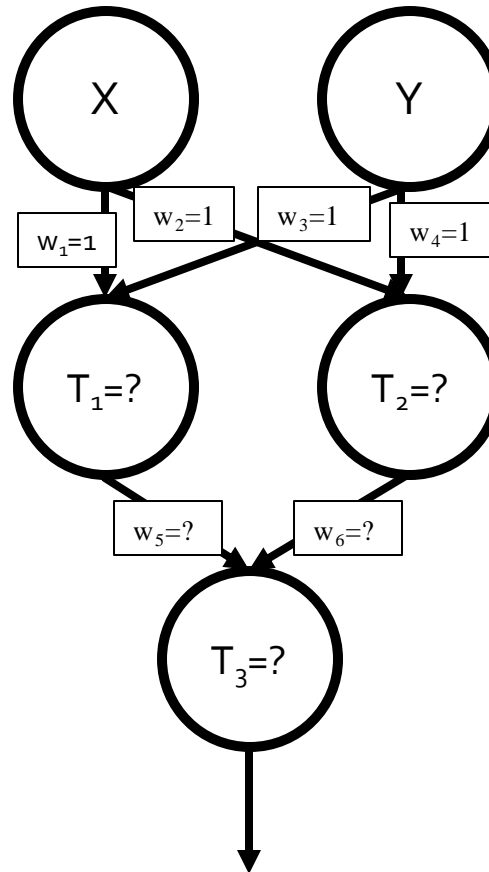
- We saw that XOR can't be solved with a single perceptron
- How about XOR with an MLP?
- Here's one example:

x1	x2	z
1	1	-1
1	0	1
0	1	1
0	0	-1



Your Turn: XOR

- Can you find weights that complete this XOR MLP?



Universal Approximation Theorem

- Two-Layer Networks are Universal Function Approximators)
 - Let F be a continuous function on a bounded subset of D -dimensional space. Then there exists a two-layer neural network F' with a finite number of hidden units that can approximate F arbitrarily well. Namely, for all x in the domain of F ,
$$|F(x) - F'(x)| < \varepsilon$$
- i.e., “two-layer networks can approximate any function”
- But we still might want more than two layers
 - Fewer neurons, time to learn, time to compute, etc.

Universal Approximation Theorem

- This is a powerful theorem, but...
 - “Just because a function can be represented does not mean it can be learned”
- Learning may require:
 - Insane complexity
 - Insane amounts of data
 - Insane computational resources
- Even if learned, the resulting network can lead to overfitting

For Next Class:

- Read:
 - Daume, chapter 10
- Good luck finishing Problem Set 4!

